

Institutional Behaviour Research Platform — Project Report

Author: Markandeya Varma **Date:** 18 August 2026 **Repositories:** `MICC` (V1) · `MICCV2` (V2) · `institutional-research` (V3, current)

1. Introduction

2. Why I am building this

3. The previous versions, and proof of work

4. The complete action plan

5. The data warehouse

6. The system architecture

7. What is completed

8. What is going on now

9. What is pending

10. What this is currently used for

11. Conclusion

Appendix A — How to verify this report

Appendix B — Key files

Appendix C — Glossary

1. Introduction

This report describes a research system I have been building since June 2026. The system studies the Indian stock market. Its purpose is to find out whether the buying and selling done by large institutions — mutual funds, foreign investors, insurance companies — can tell us anything useful about where share prices go next.

The project has gone through three versions in about eight weeks. This report covers all three. It explains what each version did, what it produced, why the first two were stopped, and what the current version is doing.

I want to say clearly at the start what the honest position is, because the rest of the report only makes sense with this in view:

The system has not found a way to make money. Two versions have been built and stopped. The current version is about 15% built. The most likely outcome, based on everything measured so far, is that the answer is "there is no usable edge here" — and the project is now deliberately designed so that it can say that clearly instead of avoiding it.

That may read like failure. I would argue it is the opposite. The first two versions could not have told me they were failing, and that is why they ran for eight weeks. The third version is built so that it can.

1.1 What the system actually does

Every day, the two Indian stock exchanges publish a list of large trades. When somebody buys or sells more than 0.5% of a company's shares in one day, the exchange must publish it: the date, the company, the name of the buyer or seller, the quantity, and the price. This is public information. Anybody can download it.

The research question is simple to state:

When a large institution buys a share, does that share do better than it otherwise would have? And if so, is the effect big enough and reliable enough to be worth acting on, after costs?

The question is simple. Getting a trustworthy answer is not, and most of this report is about why.

1.2 Reading this report

All numbers in this report were measured from the actual repositories and data on 17–18 August 2026. Nothing is estimated. Where a number is uncertain I say so. Git commit hashes are given so any claim can be checked.

2. Why I am building this

2.1 The personal reason

I want to learn how to do quantitative financial research properly. Not how to build a trading bot — that is the easy and useless part — but how to tell the difference between a real pattern and a coincidence. That distinction turns out to be the entire subject.

2.2 The intellectual reason

Financial markets produce enormous amounts of data, and almost all patterns found in that data are false. If you test ten thousand ideas against past prices, some of them will look brilliant purely by chance. This is not a small problem or a technicality. It is the central problem of the field.

The professional literature has developed real tools for this — pre-registration, multiple-testing corrections, out-of-sample testing, power analysis. Most amateur work ignores them completely, finds something that looks wonderful in a backtest, and loses money. I wanted to build something that uses the real tools.

2.3 The specific reason this dataset

Bulk and block deal disclosures are unusually good for this kind of study for three reasons:

1. **They are events with a date.** You know exactly when the information became public, so you can measure what happened afterwards without guessing.
2. **They are free and long.** India has published them since 2006. I have 223,450 bulk deals and 12,430 block deals going back twenty years.
3. **There is a plausible story.** Large institutions have research departments. If anybody has an information advantage, they should. So there is a reason to expect an effect, rather than just data-mining and hoping.

Point 3 matters more than it looks. Searching data with no idea of what you are looking for is how false discoveries are made. Having a reason in advance is what separates research from fishing.

2.4 What I am not building

I want to be exact about this, because it affects the design:

- **No live trading.** No code that places orders exists anywhere in the current repository. There is an automated test that fails the build if anybody adds functions named `place_order`, `modify_order` or `cancel_order`.
- **No money is at risk.** Nothing is connected to a real trading account.
- **No product.** This is not a startup or a service. It is a research project.

3. The previous versions, and proof of work

Three versions exist. Two are stopped. All three are in Git, so the history below can be independently verified.

3.1 Version 1 — MICC (28 June – 8 July 2026)

Repository: `~/Workspace/MICC`

Measure	Value
Commits	67
Period	28 June 2026 → 8 July 2026 (11 days)
Files	191
Python code	118 files, 23,172 lines
Final commit	<code>e27991f</code> — <i>"Portable path layer + SHP backfill completion + verification hardening"</i>

First and last commits:

2026-06-28	2af2cb1	Initial commit
2026-06-28	ce386d8	Clean data-extraction project: <code>data_extraction/</code> + <code>data_storage/</code>
2026-06-28	5adf4bd	Fix all extractors: bootstrap, table creation, dead URLs, parquet export
...		
2026-07-07	9cf8af0	Frontend: Portfolio page, idea-card lifecycle UI, <code>resilience/a11y/perf</code> + tests
2026-07-08	e27991f	Portable path layer + SHP backfill completion + verification hardening

What V1 was for. V1 was a data collection project. Its job was to gather Indian market data from public sources and store it in a usable form. It did that job, and it did it well.

What V1 produced. The lasting output of V1 is a dataset of **116 tables containing 11,276,328 rows, totalling 1.2 GB**. This still exists, and the current version is built on it. The largest tables:

Table	Rows
shp_institutional_summary	3,565,899
shp_promoter_group	2,526,543
window_extremes	1,087,140
shp_category_summary	892,410
pit_universe	359,047
features_monthly	343,963
insider_trading	283,281
bulk_deals	223,450
global_indices_daily	212,955

Why V1 stopped. It did not fail. It was replaced because it had grown into a mixture of data collection, a web frontend and analysis, all in one codebase, with no clear separation. Rather than untangle it, I started V2 and carried the data across. **This was a reasonable decision and the data has proved its value — it is irreplaceable, because the exchanges do not serve twenty years of history on request.**

3.2 Version 2 — MICCV2 (11 July – 13 August 2026)

Repository: ~/Workspace/MICCV2 · **Frozen at tag** frozen-2026-08-16

Measure	Value
Commits	107
Period	11 July 2026 → 13 August 2026 (34 days)
Files	411
Python code	204 files, 35,762 lines
Tests	45 files, 375 test functions
Reports generated	136

What V2 was for. V2 was an attempt to build a complete research and paper trading system: data warehouse, signal engines, a strategy factory that would generate and test candidate strategies automatically, a portfolio simulator, and a dashboard.

The engineering was good. This is worth stating plainly. V2 has a real data warehouse, working data collection, 375 automated tests, and a discipline of generating verification reports. As a piece of software engineering it is competent work.

The research produced nothing. This is also worth stating plainly.

The strategy factory generated and tested candidate strategies for about five weeks. The number promoted to live use was **zero**. This is not my interpretation — the phrase "0 promotions" appears in **18 separate automatically-generated verdict reports** in the repository.

The final scorecard, from `reports/p7_verdict_2026-08-15.md`:

Metric	Portfolio	Nifty 500 TR	NIFTY 50
Sharpe ratio	1.10	0.67	0.39
CAGR	8.1%	9.2%	4.6%
Volatility	7.4%	14.9%	13.7%
Max drawdown	-7.9%	-18.5%	-15.8%

The system's own pass mark was a Sharpe ratio of 1.3. It scored 1.10. Its own report calls this **"BELOW FLOOR"**.

Worse, and to the credit of whoever wrote that report, it says so itself:

"Of 648 book days, 624 are REPLAY (backtest of the champion) and 24 are inside the live-forward window. Of those 24 live-window days, only 9 were CAPTURED LIVE."

So the 1.10 Sharpe is mostly a backtest, and the live evidence is 9 days.

Why V2 was stopped. Three reasons, all found during an audit in August 2026.

Reason 1 — the cost model was wrong. V2 calculated trading costs incorrectly in four separate ways: it charged Securities Transaction Tax on one side of a trade instead of both, used the wrong exchange transaction rate, and applied GST to the wrong base. Total understatement: **10.04 basis points per round trip**.

This is not a rounding error. V2's best seasonal finding, a "turn of month" effect, showed **+3.70 basis points** of profit. With the costs corrected it becomes **-6.36 basis points** — a loss. The system's headline finding was an artefact of its own accounting mistake.

Reason 2 — it could not fail. V2 recorded 22 killed ideas and zero promotions, over about two years of accumulated development effort, and never once concluded that the whole approach might not work. There was no rule anywhere that could stop the project. A system that can only ever say "keep going" is not performing a test.

Reason 3 — documentation had drifted from reality. V2's README described an automated schedule that no longer matched the actual schedule on the machine. Nothing checked, so nobody knew.

The V2 audit is itself part of the work. Finding these three things took several days and required reading the whole system. That audit is why V3 exists and why V3 is designed the way it is.

3.3 Version 3 — institutional-research (16 August 2026 – present)

Repository: `~/Workspace/institutional-research`

Measure	Value
Commits	10
Period	16 August 2026 → present
Python code	3,316 lines
Tests	146, all passing
Decision records	18

V3 is deliberately much smaller than V2 and does much less. It is described in the rest of this report.

3.4 Summary of the three versions

	V1	V2	V3
Duration	11 days	34 days	ongoing
Commits	67	107	10
Lines of Python	23,172	35,762	3,316
Tests	—	375	146
Main output	11.3M rows of data	136 reports, 0 promotions	discipline framework
Status	superseded	frozen	active
Honest verdict	succeeded at its job	engineering good, research empty	too early to say

4. The complete action plan

The plan for V3 is written in three documents totalling about 100,000 characters (docs/plan/PLAN_1_FOUNDATIONS.md , PLAN_2_METHODODOLOGY.md , PLAN_3_EXECUTION.md). This section summarises it.

4.1 The core idea

The plan is built around one insight from the V2 audit: **the hard part is not finding patterns, it is not being fooled by them.** So the system is built discipline-first. The rules that stop self-deception were built before the research code, not after.

4.2 The eight phases

Phase	What it does	Status
0	Audit V2, freeze it, decide what to carry forward	done
0.5	Test whether data sources actually work	done
1	Foundations: paths, hashing, migrations, discipline framework	done
2	Data collection: daily automated capture of new deals	partly done
3	Identity layer: match company names to companies	not started
4	Outcomes: calculate what happened after each deal	not started
5	Benchmarks: compare against similar companies	not started
6	The four studies	not started
7	Seasonality rescan	not started
8	Reporting	not started

4.3 Two tracks

The project runs two research tracks in parallel. They share the discipline machinery — registration, decision records, multiple-testing correction — and almost nothing else, because the statistics of an event and the statistics of a calendar pattern are genuinely different.

	Track D — institutional deals	Track S — mass pattern search
Question	do disclosed institutional trades predict returns?	do <i>any</i> patterns, found by searching, survive out of sample?
Data	223,450 bulk + 12,430 block deals	21 years of prices, 4,200 stocks
Unit	one deal event	one calendar cell or signal combination
Machinery	built	not built

Track S was under-specified until 18 August. It existed in configuration files and prose but had no code and no proper design.

Worse, when the two tracks were finally checked against each other, they had never been connected. Three faults, each of which would only have appeared once scan code ran:

1. The search track had **no held-back data at all**, while the deal track sat behind a guard that refuses access.
2. A search study **could not be registered** — the checklist of required controls covered the deal track and left the search track, which has far more opportunity to fool itself, with none.
3. Two configuration files disagreed about how attempts are counted. Read literally, running the search once would have **raised the deal track's evidence bar so high that its one completed experiment would retroactively have failed** — and no future deal study could ever have passed.

All three are fixed. Attempts are now counted per research family rather than in one pool, so a search charges its own budget and not another track's; the search track has its own held-back period (2016

onward); and search studies face five required controls of their own. A fourth fault was found inside that fix — the new counters were declared permanent and nothing actually incremented them, which is the same defect the whole scheme exists to prevent, one level up.

4.3.1 Track D — the four deal studies

The deal track comes down to four questions.

Study 1 — Do institutional buys predict anything? When an institution buys, does the share beat comparable shares over the next 1, 2, 3, 5, 10 or 21 trading days?

Study 2 — Do institutional sells predict anything? The same question for selling. There is a reason to think selling is more informative: institutions buy for many reasons (money coming in, index tracking, rebalancing) but sell for fewer. 34,270 sell events have never been examined.

Study 3 — Does agreement between institutions mean more? When three or more different institutions buy the same share within 21 days, is that stronger than one institution buying? 10,098 such events exist.

4.3.2 Track S — the search track, and the 31.9 million combinations

This is half the project, and it is the half where it is easiest to fool yourself.

What the 31.9 million are. A "calendar pattern" is a question of the form: *does this share tend to rise during a particular stretch of the year?* Each one is a combination of a window length, a starting point in the calendar, and a company:

13 window lengths × 242 starting points × 4,200 companies = 13.2 million
plus index-level variants, four calendar alignments, two ways of measuring
= 31,893,556

V2 already scanned all of them, and found nothing. Its own conclusion:

"Essentially none of it survives contact with its own null. The single best pattern found — a 3-day window that rose in 94.7% of years — sits at the 94th percentile of what randomly rotated data produces, which is to say it is an ordinary result of looking 31.9 million times."

And the part that matters most: **the scan of millions found nothing, while eight carefully-reasoned guesses found two.** Both of those two were then killed by the corrected cost model.

So why do it again? Because V2 only ever asked *"is this better than chance in the data I have?"* It never asked *"does it happen again?"* — which is the only question a pattern claim actually makes. Track S asks the second one: find a pattern in 2005–2015, then require it to repeat in 2016 onward, which it has never seen.

What measuring it honestly revealed

Three findings, each of which changed the design before any code was written.

A calendar pattern happens once a year. So a company with 21 years of history gives 21 observations for any one pattern — total, ever. Split into training and testing, that leaves two to five. With two observations you could only detect an effect of about **50% a year**. Per-company calendar

analysis is not difficult; it is arithmetically impossible. And the history is not there either: the median company has **5.5 years** of prices, and only 513 of 4,200 have twenty.

Combining companies does not rescue it. The obvious fix is to pool all 4,200 together. But every company experiences a given calendar day *simultaneously*, so they all share whatever the market did that day. Measured on real data, the average pairwise correlation of daily returns is **+0.235**, which reduces 21,000 apparent observations to about **four** genuinely independent ones.

And the obvious repair turned out to be empty. Subtracting the market's own movement removes the shared part — but if "the market" is defined as the average of those same companies, then the average of what is left is **exactly zero, on every single day**. I had designed a measurement incapable of producing a number, and only found it by checking.

What survived

One approach. Instead of asking "*is this stretch of the calendar good?*", ask "**does it rank companies consistently, and does that ranking hold up in years it has never seen?**" A persistent *ordering* is far harder to produce by luck than a persistent average — chance does not put four thousand things in the same order twice.

Measured, this can detect a signal strength of about **0.014**, against typical real-world values of 0.02 to 0.05. It is the only formulation in this project that can see something real, and that single fact is why the search track is worth building at all.

The deliverable is not a list of patterns

Picking the best of 31.9 million is hopeless — the best of that many coin flips looks impressive too. So the headline result is deliberately a different thing:

Across many training-and-testing rounds, how often does a pattern chosen in training actually win in testing?

If the answer is about half the time, then searching does not work on this data — a real and useful finding that goes well beyond this project. If it is meaningfully better than half, that is a durable discovery about *method*, which is worth far more than any single pattern.

This reframing is what makes searching 31.9 million things defensible rather than reckless: **the size of the search stops being a liability and becomes the measuring instrument**. The wider you search, the more precisely you measure how badly searching overfits.

Honest expectation, stated in advance: most likely the answer is "about half the time", and almost nothing survives. That is the outcome I expect and it is still worth having.

4.4 The rules every study must follow

These are the heart of the plan.

Rule 1 — Write down the answer you will accept, before you look. Before a study runs, its hypothesis, method, and pass mark are written to a database and locked. A database trigger physically refuses to let them be changed afterwards. This is verified by trying to change one and being blocked.

Rule 2 — Exploring must be free. The data is split into three parts: 30% for free exploration, 20% for choosing between candidates, 50% for final confirmation. You may look at the exploration part as

much as you like at no cost. The confirmation part may be used once per registered study, enforced by code that raises an error.

Rule 3 — Count everything you looked at. If you test 171 ideas, the best one will look good by luck alone. The system calculates how good: with 171 attempts, pure noise produces a best result of about $t = 2.94$, so the bar is set at $t = 3.71$. With 31.9 million attempts the bar rises past $t = 7$, and for patterns measured on few observations it rises further still. These figures were themselves corrected on 18 August — see §7.6.

Rule 4 — Check the same list of confusions every time. Ten standard alternative explanations are checked automatically for every study. Any that do not apply must be dismissed in writing.

Rule 5 — An event study is not a strategy test. A finding must pass two separate gates: does the event move prices, *and* does an actual portfolio built on it beat the same portfolio without it, after costs. Section 7.3 explains why this rule exists.

Rule 6 — The project can be abandoned. If three of the four studies fail, or if nothing has passed by 28 February 2027, the conclusion is written up as a negative result and the project stops.

5. The data warehouse


5.1 Design principles

Raw data is never changed or deleted. Files are downloaded, hashed, stored compressed, and never touched again. If a file is understood wrongly, it can be re-read later. V2 stored nothing raw and had 1.2 GB of data it could not re-derive.

Every day not collected is lost forever. The exchange publishes today's deals at one web address, and replaces it tomorrow. The historical service returns an error. So collection is time-critical: about 27 trading days are already permanently missing.

Derived data can always be rebuilt. Anything computed from raw data can be deleted and regenerated. Only the raw layer is precious.

5.2 The layers



Legend. Green = working. Red = blocked. Blue = built. Yellow = designed but not yet built.

5.3 Honest status of the warehouse

Layer	Designed	Built	Note
Raw archive	yes	yes	3 files captured so far; 16 KB
V1 seed carried across	yes	no	still sits in the V2 folder, not copied
Marts	yes	no	0 of 7 tables exist as real schemas
Governance	yes	yes	9 tables, live and enforcing

The marts are the gap. Seven tables are described in the plan in detail, with column lists and constraints, but none of them has been created. They exist as prose in a document, not as a database.

5.4 Two data problems found by measurement

These were found on 17 August 2026 by running real queries. Both are serious and neither was in the plan.

Problem 1 — one third of deals cannot be matched to a price.

Of 25,097 institutional buy events, only 16,517 could be linked to price data:

Reason	Events	Share
Company not in the price database at all	5,080	20.2%
Deal happened before price coverage begins	1,808	7.2%
Deal happened after price coverage ends	83	0.3%
Price data exists but not for that exact day	1,609	6.4%
Successfully matched	16,517	65.8%

The plan sets a 5% limit for unmatched records. The reality is 34.2% — seven times over. Every result computed so far uses only the surviving two thirds.


Problem 2 — there is no historical industry classification.

The plan's main method for fair comparison requires knowing what industry a company was in at the time of the deal. The only industry data available covers **1,415 companies out of 4,200 (33.7%)**, and it is a single snapshot taken between 30 June and 2 July 2026. Applying 2026 industry labels to a 2005 event means using information that did not exist yet — the exact error the whole project is designed to avoid.


This is currently an unsolved problem, not a task. I do not have a source for it.

6. The system architecture

6.1 Overview



6.2 How a study flows through the system



6.3 Why the architecture looks like this

Each part exists because of a specific V2 failure.

Component	The V2 failure it prevents
Explicit environment (<code>paths.py</code>)	V2's checker defaulted to "dev", failed there, and rebuilt the <i>production</i> warehouse while reporting on dev
Read-only snapshots	V2 served its dashboard from the live database; 7 of 12 pages returned errors whenever anything was writing
Write-once raw archive	V2's raw folders were all zero bytes — data went straight into the warehouse and could not be re-derived
Locked experiment specs	V2 could not tell a prediction from an explanation made afterwards
Trial counter that sets bars	V2 counted attempts and never used the count for anything
Two gates	V2 had no portfolio test at all, so a finding could pass on event evidence alone
Documentation drift test	V2's README described a schedule that had changed
Provenance graph	V2 could not answer "which numbers came from this data?"

7. What is completed

7.1 The discipline framework — complete and working

This is the substantial achievement so far. 146 automated tests, all passing.

Power analysis (`power.py` , 311 lines). Answers "could this study have seen the effect even if it were there?" before running. If the answer is no, the verdict is UNDERPOWERED — reported as *silence*, not as evidence of absence. This distinction did not exist in V2.

Data split (`split.py` , 237 lines). Divides companies into three groups by a mathematical hash of their ISIN code. Exploration is free; confirmation is enforced by code that raises an error.

The choice of ISIN rather than company symbol was driven by measurement:

Finding	Value
ISINs with more than one symbol over time	276
Those symbols appearing in deal data	459
Deal records affected	26,046 (11.04%)

Real examples: `CADILAH` became `ZYDUSLIFE` ; `PRISMCEM` became `PRSMJOHNSN` . If the split used symbols, those companies would sit in the exploration group under one name and the confirmation group under the other — contaminating the confirmation set for 11% of the data, with nothing looking wrong in any output.

Multiple testing (`multiplicity.py` , 186 lines). Converts "how many things did we try" into a required standard of evidence:

Attempts	Best result from pure noise	Required standard
10	t = 1.90	t = 2.80
171	t = 2.94	t = 3.71
1,000	t = 3.45	t = 4.31
31,893,556	t = 5.63	t = 7.04

(Corrected 18 August. The figures first published were lower — see §7.6.)

Design gate (design.py , 275 lines). Refuses to let a study exist unless it has a mechanism, a prediction that could fail, computed detectability, and a plan for all ten standard confusions.

The test that proves this is worth having is that **it rejects the project's own first experiment**. More on that in 7.3.

7.2 Documentation that cannot silently rot

Twenty-six decision records, each recording what was decided, by whom, why, **what would reverse it**, and what it costs. An automated test fails the build if any record is missing those fields.

A second test compares the plan documents against the configuration files. On its first run it found three real errors: the plan still advertised abandoned time horizons, never mentioned the portfolio gate, and never mentioned the data split. All three were fixed.

7.3 A completed experiment, and why it matters

One full experiment has been run start to finish: exp_001 .

The hypothesis. A portfolio of the largest 500 Indian companies which *avoids* any company that had an institutional buy in the last 10 trading days will do better than the same portfolio without that rule.

The sequence, which is the point. The prediction and the pass mark were written to the database and locked in commit f25608d . The test was run afterwards. The result was recorded in commit c31e128 . **The registration is earlier in the Git history than the result**, so it cannot have been written to fit the answer.

What happened at the event level. The effect was real and survived three attempts to kill it:

Check	Result
vs random companies, same dates	−0.860% specific to the event
vs companies of similar volatility	−0.805%, t = −3.93
Is it just price reversal?	No — correlation +0.008, pattern not monotonic

What happened at the portfolio level.

Half of the data	Annual difference	t	Verdict
First half	+0.237%	+3.11	FAIL
Held-back half	−0.022%	−0.25	FAIL

Why it died. The rule only excluded **1.2% of companies** at any time. A -0.8% effect applied to 1.2% of a portfolio is about one basis point a month — nothing.

This is the most valuable thing learned so far. Every statistic in the event study was correct and pointed the same way. The study was still useless. Only building the actual portfolio revealed it. **An event study is not a strategy test**, and the gap between them is not a technicality.

Under V2's design this would have been recorded as a success.

7.4 A measurement that corrected me

On 17 August I stated that all the project's detectability calculations were too optimistic because they treated companies as independent. **This was wrong**, and I found it wrong within hours by measuring it:

Method	Detectable effect
Treating 16,445 events as independent	0.076% ← never actually done
What the code actually does (monthly grouping)	0.621%
With the correction that <i>was</i> missing	0.660%

The eight-fold penalty was already being applied. The real gap was a different and smaller one — correlation between months, worth 6–27%, not a factor of three.

This is recorded in decision record 0017 rather than quietly edited out. I include it here because a report that only lists successes is not a report.

7.5 Daily data capture — running

A collector runs automatically at 20:00, 22:30 and 08:00 IST. On its first night it captured 100 bulk deals and 4 block deals, verified by hash. Files are compressed, hashed, never overwritten, and duplicates are detected.

7.6 Four errors of my own, found by verification

On 18 August I was asked simply "*are you sure about the whole thing?*" Checking found four errors in work committed within the previous day. I record them because a report listing only successes is not a report, and because how they were found matters more than that they existed.

Error 1 — a statistic that measured itself. I reported that removing the market's overall movement from returns made stocks statistically independent, and built a design on it. In fact the calculation I used *forces* that answer no matter what the data contains. Feeding it a strongly-linked dataset and a completely random one produced the same result. It measured the arithmetic of the operation, not the market.

Error 2 — a question with no possible answer. Worse, the way I had defined "market-relative" meant that the quantity I proposed to study is exactly zero for every day by construction. I had designed a measurement incapable of producing a number. The correct approach — comparing *rankings* of stocks rather than their average — was in the plan as one option among several. It is in fact the only one.

Errors 3 and 4 — the significance bar was too low, twice. The tool that sets how strong evidence must be had two independent faults, and both made results *easier* to pass:

Fault	Effect
Measured one direction of movement while being applied to both	every bar understated
Assumed a bell curve where the true distribution has heavier tails	a further 23% understatement for small samples

Corrected, the bar for the deal track rises from $t = 3.62$ to **$t = 3.71$** . The one completed experiment still clears it, so **no recorded verdict changes**.

Why they survived 146 tests. The tests compared the tool against *its own output*, pinned as expected values. A test asserting that code agrees with itself proves nothing. The replacement tests compare against independent simulation, which is the standard the originals should have met.

What this says about the project. The framework caught none of these — a person asking a direct question did. The safeguards are real and they are not sufficient, and I would rather state that than present a system that appears to police itself.

8. What is going on now

8.1 Measuring whether there is anything to find

The current activity is establishing what size of effect this data can possibly reveal. Results from 17 August, on 16,445 real institutional buy events:

Horizon	Observed effect	Smallest detectable	Conclusion
1 session	+0.691%	0.191%	detectable
5 sessions	-0.010%	0.423%	nothing
10 sessions	-0.603%	0.660%	below the floor
21 sessions	-1.061%	0.968%	detectable but implausibly large

Two things stand out.

The sign flips. Shares rise 0.691% on the first day, then fall, ending 1.061% down after 21 days. This is the classic signature of a temporary price push that then reverses — consistent with the finding that **54.8% of bulk deals are the same party buying and selling on the same day**, which is market-making activity, not investment.

The earlier finding weakens. The 10-session effect of -0.603% is *smaller* than the 0.660% we can reliably detect. The -0.805% found in `exp_001` came from comparing against companies of similar volatility, and it was that tighter comparison that provided the sensitivity — not the effect being robust to any comparison.

8.2 An unresolved question that affects everything

The system judges whether a study is worth running by comparing what it can detect against what size of effect is realistic. That threshold is currently a single number — 0.5% per month.

The problem: this number is compared against every time horizon, but it is only dimensionally correct at about 21 sessions. Comparing a one-day result against a monthly threshold is a units error.

The fix depends on an unanswered question about how the effect should behave:

- **If a disclosure causes a single one-off price adjustment**, a fixed threshold is right and the table above stands.
- **If institutions have persistent skill that accrues over time**, the threshold must scale with the horizon — and then **every row in the table becomes undetectable**, including the two marked otherwise.

This is recorded as open decision 0018. **Until it is resolved, no result in this project should be quoted as final.** It is the single most important open item.

9. What is pending

9.1 Build status, measured

Component	Status
Environment, paths, hashing, migrations	built
Discipline framework	built
Daily capture (basic)	built
Trading calendar	not built
V1 data copied into the project	not done
Reconciliation check	never run
Full collector	not built
Company identity matching	not built
Institution name normalisation	not built
Outcome calculation	not built
Cost model in code	config only
Benchmark construction	config only
The four deal studies (Track D)	not built
Track S — walk-forward folds	not built
Track S — calendar scan	not built
Track S — signal combinations	not built
Track S — the procedure test	not built
Provenance graph populated	0 rows

Seven of seven core data tables do not exist, and all seven Track S modules are unwritten. The project folder currently holds 16 KB of data.

About 15% of the system is built, and the completed portion is the framework rather than the research machinery.

9.2 Three things in the plan that cannot be built as written

1. No historical industry data (Section 5.4). The plan's main comparison method needs it. It covers 33.7% of companies and is a single 2026 snapshot. There is no known source. **This is a genuine blocker, not a task.**

2. Company matching is unspecified. The plan sets a 5% limit and gives no method for reaching it. Measured reality is 34.2%.

3. The seasonality rescan has no time estimate. "About three weeks" is an assertion; no single unit of the calculation has ever been timed, and a later decision roughly quadrupled the work.

9.3 Risks

Risk	Status
No backup of any kind	open — no remote repository, no external drive. Captured trading days cannot be re-obtained at any price
Data lost each uncollected day	reduced — collector now runs, but ~27 days already lost
BSE data unavailable	open — all routes return errors
Historical backfill impossible	open — service returns error 503
No detectable effect exists	the most likely outcome

9.4 Timeline, and why it had to be rewritten

The plan originally listed nine phases in sequence and stated that no deadline had been set. **That was false** — a deadline of 28 February 2027 had been fixed two days earlier. Measured honestly, at the 2–3 hours a day I actually have:

	Weeks
calendar available	27.9
work planned	26.2
slack	+1.7

Six percent. On a project where **every estimate so far has been wrong**, a 25% overrun — which is optimistic for software — missed the deadline by five weeks.

The fix is not a better estimate. It is deciding **in advance** what gets cut.

The critical path — 14 weeks to one defensible answer

Phase	Weeks	Ends
Warehouse and reconciliation	3	7 Sep 2026
Collection (<i>reduced — the daily capture already runs</i>)	1	14 Sep 2026
Identity layer (<i>where the 34.2% matching failure lives</i>)	4	12 Oct 2026
Clean data mart	2	26 Oct 2026
Costs and benchmarks	2	9 Nov 2026
One outcome study	2	23 Nov 2026

This lands one week inside the checkpoint, with **13.9 weeks of buffer** — 36% overrun tolerance in place of 6%.

Two reductions make it fit, and neither is guesswork. Collection drops from three weeks to one because the daily capture has been running since 17 August and the raw bytes are already safe; the full parser can wait for a study that needs it. The outcome study drops from four studies to **one**, because one study answered is worth more than four half-answered.

What gets cut, decided now

Extension	Weeks	Cut order
Re-run the first experiment reproducibly	0.2	last
Studies 2–4	3	
The search track	4	
Seasonality	1	
Monitoring and automated reports	2	first

Under deadline pressure, what gets dropped is whatever is least *defended* at that moment — not whatever matters least. Deciding the order now, while nothing is at stake, is the only point at which that decision is honest.

The seasonality decision changed as part of this. I had chosen to rebuild all 31.9 million calendar cells from scratch — three weeks — against advice to validate the existing ones instead. Faced with the schedule arithmetic I reversed that: validation takes one week, and those two weeks are what the buffer is made of. Validation is not "trust the old numbers": a 100,000-cell sample is recomputed from scratch and must match **exactly**, any mismatch escalates automatically back to a full rebuild, and every statistical conclusion is recomputed regardless. What is reused is the raw cell counts, which are cheap to check and expensive to regenerate.

Honest limit. This makes the deadline survivable for the thing that matters — one answer by the checkpoint — but not for everything. Running the full extension list still misses a 25% overrun by 2.4 weeks. That is what the cut order is for.

Date	Milestone
23 November 2026	Critical path complete — first portfolio-gated verdict
30 November 2026	Checkpoint — if no study has reached even the first gate, stop early
28 February 2027	Deadline — if nothing has passed, write the negative result and stop

10. What this is currently used for

I want to be accurate here rather than flattering.

10.1 What it is used for today

Nothing financial. No money is invested. No orders are placed. No recommendations are produced.

The system currently has three real uses:

1. **Daily data capture.** It is preserving Indian institutional deal data that would otherwise be permanently lost. This has value regardless of whether the research succeeds.
2. **A method for testing ideas honestly.** The framework — pre-registration, split data, multiple-testing correction, two gates — is general. It applies to any field where you search data for patterns.
3. **Learning.** I now understand why most quantitative research fails, because I have made the mistakes and measured them.

10.2 What it is not used for

- Not trading, and not connected to any account
- Not advising anyone
- Not producing predictions
- Not producing any validated finding — **zero results have passed both gates**

10.3 What it could be used for if it works

If a study passes both gates and survives the deadline, the output would be a documented, reproducible finding about Indian institutional behaviour, with its provenance recorded. Whether that would be worth trading is a separate question requiring capital, execution and risk management that do not exist here.

Based on the evidence so far I would put the probability of a tradable finding below 20%.

The signals point the wrong way: the effect reverses across horizons, most bulk-deal activity is market-making rather than investment, the one completed experiment failed, and the detectability question in Section 8.2 may render the whole range unusable.

10.4 The most likely honest outcome

The most likely result is a well-documented negative: *public institutional deal disclosure in India does not contain a tradable edge for a retail participant after realistic costs.*

That is a legitimate result. It is also unpublishable in most amateur settings, which is precisely why so much amateur work reports positives.

11. Conclusion

11.1 What has been achieved

In roughly eight weeks I have built three versions of a market research system, audited two of them to destruction, and produced:

- A dataset of 11.3 million rows spanning 2005–2026, still in use
- A working data warehouse and collection system (V2), now frozen
- A discipline framework with 237 tests that enforces honest research
- One complete experiment, correctly rejected by its own pre-registered rule
- Twenty-six decision records with reversal conditions
- Four material measurements that changed the plan: a 10.04 bp cost error, 54.8% market-making contamination, a 34.2% matching failure, and a missing industry history

11.2 What has not been achieved

- **No validated finding.** Zero results have passed both gates.
- **No profitable strategy.** V2 promoted zero strategies in five weeks.
- **The V3 system is ~15% built.** No core data table exists.
- **Two plan components have no data source.**
- **One critical question is unresolved** and may invalidate every current result.

11.3 The honest assessment

If this project is judged on "did you find a way to beat the market", the answer is no, and probably will remain no.

If it is judged on "did you build something that can tell the difference between a real finding and a lucky one", the answer is yes, and that is the harder problem.

The difference between V2 and V3 is not sophistication — V2 has ten times the code. It is that V2 could not fail and V3 can. V2 ran for five weeks producing zero promotions and never concluded anything was wrong. V3 has a written rule that stops the project on 28 February 2027 if nothing works, and a first experiment that was killed by its own pre-registered standard within ninety minutes of looking promising.

11.4 What I have learned

Most patterns in financial data are false. With 171 attempts, noise alone produces a t-statistic of 2.94 — which most people would report as highly significant.

Being able to fail is a feature. V2's real defect was not any single bug. It was that no possible result would have caused it to stop.

Statistics that are individually correct can be collectively useless. exp_001's event study was correct in every particular and the strategy was worthless, because the rule touched 1.2% of the portfolio.

Measurement beats assumption, including about yourself. Four significant errors were found by measuring rather than reasoning — including one of my own, made and corrected within hours on 17

August.

11.5 Next steps

These are the critical path of §9.4, in order. Everything else is an extension that gets cut before this list does.

1. **Answer the detectability question** — one judgement, and it gates everything downstream. The code now *refuses* to register the affected studies rather than guessing, so this is a hard block rather than a nagging one.
2. **Arrange a backup.** A single disk failure currently loses irreplaceable data, including trading days that cannot be re-obtained at any price. Cheap, and it is the only item here whose cost of neglect is permanent.
3. **Copy the data in and run the reconciliation check.** Nothing downstream can be trusted until it passes, and it is also the cheapest way to find out whether the 34.2% matching failure is fixable or structural.
4. **Establish whether the industry-history problem is solvable.** If it is not, the primary comparison method has to change, and I would rather know that in September than in January.
5. **Then run one study, properly** — institutional selling, 34,270 events, never examined. One study through both gates by 23 November is the deliverable. Studies 2 to 4 are extensions.

Appendix A — How to verify this report

```
# Version 1
cd ~/Workspace/MICC && git rev-list --count HEAD          # 67
git log --reverse --format='%ad %h %s' --date=short | head -1

# Version 2
cd ~/Workspace/MICCV2 && git rev-list --count HEAD        # 107
git tag                                                    # frozen-2026-08-16
grep -l "0 promotions" reports/*.md | wc -l              # 18
sed -n '1,30p' reports/p7_verdict_2026-08-15.md

# Version 3
cd ~/Workspace/institutional-research
RESEARCH_ENV=dev python -m pytest tests/ -q              # 146 passed
git log --oneline
ls docs/decisions/                                       # 18 records

# The pre-registration ordering – registration BEFORE result
git log --oneline --reverse | grep -E "f25608d|c31e128"
```


Appendix B — Key files

File	Purpose
src/research/power.py	Can this study see anything?
src/research/split.py	Explore / select / confirm partition
src/research/multiplicity.py	How much did we look?
src/research/design.py	Is this study allowed to exist?
src/archive/stopgap.py	Daily capture
configs/*.yaml	All parameters — single source of truth
docs/decisions/	18 decision records
docs/plan/PLAN_[1-3]_*.md	The full plan

Appendix C — Glossary

Term	Meaning
Bulk deal	A trade above 0.5% of a company's shares in one day; must be disclosed
Block deal	A large pre-arranged trade in a separate exchange window
MDE	Minimum detectable effect — the smallest effect a study could reliably find
t-statistic	How many standard errors a result is from zero; above ~2 is conventionally "significant"
Sharpe ratio	Return divided by volatility; higher is better
Basis point	One hundredth of a percent
Pre-registration	Writing down predictions and pass marks before looking
ISIN	A permanent 12-character code identifying a company's shares. Unlike a ticker symbol it does not change when a company is renamed
Monotonic	Moving in one direction only — a counter that can rise but never fall
Point-in-time	Using only information that existed on the date in question
Provenance	A record of which data and code produced a number
UNDERPOWERED	The study could not have detected the effect either way — silence, not evidence